

MATLAB **EXPO**



Reinforcement Learning Workflows for AI




Key Takeaways

- What is reinforcement learning and why should I care about it?
- How do I set up and solve a reinforcement learning problem?
- What are some common challenges?


Why Should You Care About Reinforcement Learning?








One Approach Could Be...



Any Alternatives?





Applications of Reinforcement Learning





What is reinforcement learning?

Type of machine learning that trains an '**agent**' through trial & error interactions with an **environment**


Reinforcement Learning vs Machine Learning vs Deep Learning




Reinforcement Learning vs Machine Learning vs Deep Learning




Reinforcement Learning vs Machine Learning vs Deep Learning



Reinforcement Learning vs Machine Learning vs Deep Learning



Reinforcement Learning vs Machine Learning vs Deep Learning





What about deep learning?

Complex reinforcement learning problems typically need deep neural networks
[Deep Reinforcement Learning]

How does reinforcement learning training work?


Analogies with pet training



13

Reinforcement Learning Concepts


Training a self-driving car



- Vehicle's computer...
(agent)
- is reading sensor measurements from LIDAR, cameras,...
(observations)
- that represent road conditions, vehicle position,...
(environment)
- and generates steering, braking, throttle commands,...
(action)
- based on an internal state-to-action mapping...
(policy)
- that tries to optimize, e.g., lap time & fuel efficiency...
(reward).
- The policy is updated through repeated trial-and-error by a
reinforcement learning algorithm

Reinforcement Learning Concepts


Training a self-driving car




- Vehicle's computer uses the final state-to-action mapping... (**policy**)
- to generate steering, braking, throttle commands,... (**action**)
- based on sensor readings from LIDAR, cameras,... (**observations**)
- that represent road conditions, vehicle position,... (**environment**).

By definition, this trained policy is optimizing lap time & fuel efficiency

Reinforcement Learning vs Controls



- Adaptation mechanism
- Error/Cost function
- Manipulated variable
- Measurement
- Plant
- Controller



- RL Algorithm
- Reward
- Action
- Observation
- Environment
- Policy


Reinforcement learning has parallels to **control system design**

Policy Representation and Deep Learning

Representation options

- Look-up table
- Polynomials

Observations



Next action

Look-up tables **do not scale** well

Policy Representation and Deep Learning

Representation options

- Look-up table
- Polynomials
- **(Deep) neural networks**



Deep neural network policy

Observations
(camera frame, sensors, ...)



Next action

Neural networks allow representation of **complex policies**

How do I set up and solve a reinforcement learning problem?

Reinforcement Learning Workflow

- Simulation models or real hardware
- Virtual models are safer and cheaper



Environment

Reward

Policy representation



Agent

- Deep network? Table? Polynomial?

- Select training algorithm
- Tune hyperparameters



Training

- Trained policy is a standalone function




Deployment

- Numerical value that evaluates goodness of policy
- Reward shaping can be challenging

- Large number of simulations needed
- Parallel & GPU computing can speed up training
- Training could still take hours or days


Reinforcement Learning Workflow



Reinforcement Learning Toolbox

Introduced in R2019a

- Built-in and custom reinforcement learning algorithms
- Environment modeling in MATLAB and Simulink
 - Existing scripts and models can be reused
- Deep Learning Toolbox support for representing policies
- Training acceleration with Parallel Computing Toolbox and MATLAB Parallel Server
- Deployment of trained policies with GPU Coder and MATLAB Coder
- Reference examples for getting started




Reinforcement Learning Toolbox™ provides functions and blocks for training policies using reinforcement learning algorithms including DQN, A2C, and DDPG. You can use these policies to implement controllers and decision-making algorithms for complex systems such as robots and autonomous systems. You can implement the policies using deep neural networks, polynomials, or look-up tables.

The toolbox lets you train policies by enabling them to interact with environments represented by MATLAB® or Simulink® models. You can evaluate algorithms, experiment with hyperparameter settings, and monitor training progress. To improve training performance, you can run simulations in parallel on the cloud, computer clusters, and GPUs (with Parallel Computing Toolbox™ and MATLAB Parallel Server™).

Through the ONNX™ model format, existing policies can be imported from deep learning frameworks such as TensorFlow™ Keras and PyTorch (with Deep Learning Toolbox™). You can generate optimized C, C++, and CUDA code to deploy trained policies on microcontrollers and GPUs.


The toolbox includes reference examples for using reinforcement learning to design controllers for robotics and automated driving applications.



Example: Walking Robot




**Control objective: Walk
on a straight line**



Creating the Environment




rlWalkingBipedRobot_Template ▶ Walking Robot ▶





Reward Shaping



Reward function inspired by
"Emergence of Locomotion
Behaviours in Rich Environments"
Google DeepMind, 2017
<https://arxiv.org/pdf/1707.02286.pdf>



DESIGNER



LAYER LIBRARY

Filter layers...

INPUT

imageInputLayer

image3dInputLayer

sequenceInputLayer

roiInputLayer

CONVOLUTION AND FULLY CONNECTED

convolution2dLayer

convolution3dLayer

groupedConvolution2dLayer

transposedConv2dLayer



transposedConv3dLayer

fullyConnectedLayer

SEQUENCE

lstmLayer

hilstm aver



PROPERTIES

Number of layers 7

Number of connections 6

Input type Image

Output type None


Creating the Agent



Walking Robot: Reinforcement Learning

Copyright 2019 The MathWorks, Inc.



```
agentOptions = rlDDPGAgentOptions;
agentOptions.SampleTime = Ts;
agentOptions.DiscountFactor = 0.99;
agentOptions.MiniBatchSize = 128;
agentOptions.ExperienceBufferLength = 1e6;
agent = rlDDPGAgent(actor,critic,agentOptions);
```



Training the Agent




```
trainOpts.UseParallel = true;  
trainOpts.ParallelizationOptions.Mode = 'async';
```




Mechanics Explorers - Mechanics Explorer-rlWalkingBipedRobot_Template


File Explorer Simulation View Tools Window Help



□ X




Mechanics Explorer-rlWalkingBipedRobot_Template




T = [0,20] □ 1X — Time 11.137056

Applications of Reinforcement Learning





Autonomous Driving Example




Objective: Augment traditional controller with reinforcement learning to improve lap time

RL Training Environment





Results



Traditional controller + reinforcement learning



Reference Applications in Documentation

- Controller Design
- Robotic Locomotion
- Lane Keep Assist
- Adaptive Cruise Control
- Imitation Learning

The image displays six reference applications for reinforcement learning, each with a title, a brief description, and a screenshot or diagram.

- Train DDPG Agent to Control Flying Robot**
Train a reinforcement learning agent to control a flying robot.
- Train Biped Robot to Walk Using DDPG Agent**
Train a biped robot to walk using a DDPG agent. The diagram shows joints labeled Hip joint, Knee joint, and Ankle joint.
- Ego Car with Adaptive Cruise Control (ACC)**
Train DDPG Agent for Adaptive Cruise Control. The diagram shows a car's sensor fusion and control architecture.
- Train DQN Agent for Lane Keeping Assist**
Train a reinforcement learning agent for a lane keeping assist application. The diagram shows sensor dynamics and vehicle dynamics.
- Train DDPG Agent for Path Following Control**
Train a reinforcement learning agent for a lane following application. The screenshot shows plots of steering angle and linear velocity over time.
- Imitation Learning**
A placeholder card for imitation learning.

Pros & Cons of Reinforcement Learning

Pros

- **No data** required before training
- **New possibilities** with AI for hard-to-solve problems
- Complex **end-to-end** solutions can be developed
- **Uncertain, nonlinear** environments can be used

Cons

- Trained policies are **hard to verify** (no performance guarantees)
- Many trials/data points required (**sample inefficient**)
 - Training with real hardware can be expensive and dangerous
- Large number of **design parameters**
 - Reward signal
 - Network architectures
 - Training Hyperparameters

Simulations are key in Reinforcement Learning

How Can MATLAB and Simulink Help?

Challenges

- Trained policies are **hard to verify** (no performance guarantees)
- Many trials/data points required (**sample inefficient**)
 - Training with real hardware can be expensive and dangerous
- Large number of **design parameters**
 - Reward signal
 - Network architectures
 - Training Hyperparameters

MATLAB®&SIMULINK®

- **Reuse** existing code and models for environments
- Use simulations for **policy verification**
 - Simulate extreme scenarios
- Run simulation trials **in parallel** to accelerate training
- Consult Reinforcement Learning Toolbox **examples**
 - Iterative tuning with simulations

Key Takeaways

- What is reinforcement learning and why should I care about it?
- How do I set up and solve a reinforcement learning problem?
- What are some common challenges?

Learn More

- Reference examples for controls, robotics, and autonomous system applications
- Documentation written for engineers and domain experts
- Tech Talk video series on Reinforcement Learning concepts
- Reinforcement Learning ebooks available at [mathworks.com](https://www.mathworks.com)

